

11 - Neural Networks

Part 1 - XOR

1. Using the XOR dataset below, train (400 epochs) a neural network (NN) using 1, 2, 3, 4, and 5 hidden layers (where each layer has only 2 neurons). For each n layers, store the resulting loss score along with n. Plot the results to find what the optimal number of layers is.
2. Repeat the above with 3-neuron and 4-neuron Hidden layers. How do these results compare to the 2 neuron layers?
3. Using the most optimal configuraion (n-layers, k-neurons per layer), compare how `tanh` , `sigmoid` , `softplus` and `relu` effect the loss after 400 epochs. Try other Activation functions as well (<https://keras.io/activations/>)
4. Instead of `SGD` try other optimizers and report on the loss score. (<https://keras.io/optimizers/>)

```
In [1]: from keras.models import Sequential
        from keras.layers import Dense
        from keras.optimizers import SGD #Stochastic Gradient Descent

        import pandas as pd
        import numpy as np
        # fix random seed for reproducibility
        np.random.seed(7)

        import matplotlib.pyplot as plt
        %matplotlib inline
```


Using TensorFlow backend.

```
In [2]: n = 40
        xx = np.random.random((n,1))
        yy = np.random.random((n,1))
```

```
In [3]: X = np.array([np.array([xx, -xx, -xx, xx]), np.array([yy, -yy, yy, -yy])]).reshape(2, 4*n)
        y = np.array([np.ones([2*n]), np.zeros([2*n])]).reshape(4*n)
```

```
In [4]: plt.scatter(*zip(*X), c=y)
```

```
Out[4]: <matplotlib.collections.PathCollection at 0x170b56a8b08>
```



1 layer, 2 neurons

```
In [5]: model1 = Sequential()

model1.add(Dense(2, input_dim=2, activation='tanh')) #first layer
model1.add(Dense(1, activation='sigmoid'))

sgd = SGD(lr=0.1)
model1.compile(loss='binary_crossentropy', optimizer='sgd')

model1.fit(X, y, batch_size=2, epochs=400) #160/4 = 40 per epoch. we fit the data
print(model1.predict_proba(X).reshape(4*n))

scores = model1.evaluate(X, y) # evaluate the model

80/80 [=====] - 0s 1ms/step - loss: 0.6875
Epoch 36/400
80/80 [=====] - 0s 947us/step - loss: 0.6872
Epoch 37/400
80/80 [=====] - 0s 1ms/step - loss: 0.6868
Epoch 38/400
80/80 [=====] - 0s 1ms/step - loss: 0.6864
Epoch 39/400
80/80 [=====] - 0s 898us/step - loss: 0.6858
Epoch 40/400
80/80 [=====] - 0s 1ms/step - loss: 0.6855
Epoch 41/400
80/80 [=====] - 0s 972us/step - loss: 0.6851
Epoch 42/400
80/80 [=====] - 0s 972us/step - loss: 0.6845
Epoch 43/400
80/80 [=====] - 0s 1ms/step - loss: 0.6838
Epoch 44/400
80/80 [=====] - 0s 960us/step - loss: 0.6833
Epoch 45/400
```

```
In [6]: scores1 = model1.evaluate(X, y)
scores1, model1.metrics_names

5/5 [=====] - 0s 1ms/step - loss: 0.3173
```

```
Out[6]: (0.3173014521598816, ['loss'])
```

```
plt.scatter(zip(X), c=model1.predict(X).reshape(4*n))
```

2 layers, 2 neurons

```

In [7]: model2 = Sequential()

model2.add(Dense(2, input_dim=2, activation='tanh')) #first layer
model2.add(Dense(2, activation='tanh')) #second layer
model2.add(Dense(1, activation='sigmoid'))

sgd = SGD(lr=0.1)
model2.compile(loss='binary_crossentropy', optimizer='sgd')

model2.fit(X, y, batch_size=2, epochs=400)
print(model2.predict_proba(X).reshape(4*n))

scores2 = model2.evaluate(X, y) # evaluate the model

```


0.24734333 0.9337923 0.94307888 0.9389123 0.93782033 0.93393838
0.74615055 0.90717804 0.10576588 0.9549526 0.9662323 0.7780062
0.9678076 0.9685021 0.09452382 0.8860277 0.96701056 0.15198445
0.96692574 0.96899426 0.96730506 0.92001367 0.9433472 0.5926153
0.70934296 0.92005116 0.9503858 0.962932 0.9061436 0.67016655
0.90433824 0.5030033 0.18561235 0.201395 0.08401108 0.12929755
0.34707493 0.07295185 0.05253154 0.8838269 0.05291629 0.09291396
0.08042651 0.1475657 0.09748438 0.29919767 0.0612832 0.90606844
0.0556474 0.09977868 0.17404866 0.7151205 0.17144898 0.08196738
0.06314722 0.07592601 0.10373479 0.05765411 0.28540272 0.07764047
0.07306796 0.051173 0.6616037 0.1381262 0.05158785 0.05084795
0.19007108 0.12680581 0.6015444 0.05204031 0.07836878 0.05181479
0.16408855 0.23152715 0.0761663 0.1504196 0.3839514 0.06720281
0.05546361 0.8802028 0.05568403 0.08379501 0.09201756 0.1717977
0.0873971 0.27158833 0.05757689 0.911978 0.05952567 0.08952773
0.20236537 0.7062851 0.15367556 0.09465659 0.06941399 0.08612427
0.1210503 0.05470872 0.26152593 0.07144606 0.06751639 0.05136165
0.650609 0.15948981 0.05366191 0.05156997 0.16964638 0.14803526
0.6234674 0.05142802 0.08920226 0.05132195]
5/5 [=====] - 0s 1ms/step - loss: 0.3159

```

In [8]: plt.scatter(*zip(*X), c=model2.predict(X).reshape(4*n))

```

Out[8]: <matplotlib.collections.PathCollection at 0x170b6dcd9c8>



3 layers, 2 neurons

```

In [9]: model3 = Sequential()

model3.add(Dense(2, input_dim=2, activation='tanh')) #first layer
model3.add(Dense(2, activation='tanh')) #second layer
model3.add(Dense(2, activation='tanh')) #third layer
model3.add(Dense(1, activation='sigmoid'))

sgd = SGD(lr=0.1)
model3.compile(loss='binary_crossentropy', optimizer='sgd')

model3.fit(X, y, batch_size=2, epochs=400)
print(model3.predict_proba(X).reshape(4*n))

scores3 = model3.evaluate(X, y) # evaluate the model

plt.scatter(*zip(*X), c=model3.predict(X).reshape(4*n))

```

```

Epoch 45/400
80/80 [=====] - 0s 1ms/step - loss: 0.4727
Epoch 46/400
80/80 [=====] - 0s 2ms/step - loss: 0.4695
Epoch 47/400
80/80 [=====] - 0s 1ms/step - loss: 0.4677
Epoch 48/400
80/80 [=====] - 0s 1ms/step - loss: 0.4653

Epoch 49/400
80/80 [=====] - 0s 1ms/step - loss: 0.4634
Epoch 50/400
80/80 [=====] - 0s 1ms/step - loss: 0.4619
Epoch 51/400
80/80 [=====] - 0s 1ms/step - loss: 0.4597
Epoch 52/400
80/80 [=====] - 0s 1ms/step - loss: 0.4581
Epoch 53/400
80/80 [=====] - 0s 2ms/step - loss: 0.4567
Epoch 54/400

```

4 layers, 2 neurons

```

In [12]: model4 = Sequential()

model4.add(Dense(2, input_dim=2, activation='tanh')) #first layer
model4.add(Dense(2, activation='tanh')) #second layer
model4.add(Dense(2, activation='tanh')) #third layer
model4.add(Dense(2, activation='tanh')) #fourth layer
model4.add(Dense(1, activation='sigmoid'))

sgd = SGD(lr=0.1)
model4.compile(loss='binary_crossentropy', optimizer='sgd')


model4.fit(X, y, batch_size=2, epochs=400)
print(model4.predict_proba(X).reshape(4*n))

scores4 = model4.evaluate(X, y) # evaluate the model

plt.scatter(*zip(*X), c=model4.predict(X).reshape(4*n))
0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000
0.5551992 0.04853067 0.06210914 0.04820016]
5/5 [=====] - 0s 1ms/step - loss: 0.3103

```

Out[12]: <matplotlib.collections.PathCollection at 0x170b93952c8>



5 layers, 2 neurons

```

In [14]: model5 = Sequential()

model5.add(Dense(2, input_dim=2, activation='tanh')) #first layer
model5.add(Dense(2, activation='tanh')) #second layer
model5.add(Dense(2, activation='tanh')) #third layer
model5.add(Dense(2, activation='tanh')) #fourth layer
model5.add(Dense(2, activation='tanh')) #fifth layer
model5.add(Dense(1, activation='sigmoid'))

sgd = SGD(lr=0.1)
model5.compile(loss='binary_crossentropy', optimizer='sgd')

model5.fit(X, y, batch_size=2, epochs=400)
print(model5.predict_proba(X).reshape(4*n))


scores5 = model5.evaluate(X, y) # evaluate the model

plt.scatter(*zip(*X), c=model5.predict(X).reshape(4*n))

0.011390851 0.01134229 0.01162797 0.01133832]
5/5 [=====] - 0s 1ms/step - loss: 0.0482

```

Out[14]: <matplotlib.collections.PathCollection at 0x170b935bb08>



Plotting the results

```


In [15]: df = np.array([scores1, scores2, scores3, scores4, scores5])
df

```

Out[15]: array([0.31730145, 0.31588522, 0.36067224, 0.31026211, 0.0482264])

```
In [16]: plt.plot(df)
```

```
Out[16]: [<matplotlib.lines.Line2D at 0x170b9537208>]
```



Part 1 - continued

2. Repeat the above with 3-neuron and 4-neuron Hidden layers. How do these results compare to the 2 neuron layers?

3-neuron models


```

In [18]: # 1 Layer
model1_3 = Sequential()
model1_3.add(Dense(3, input_dim=2, activation='tanh')) #first Layer
model1_3.add(Dense(1, activation='sigmoid'))
sgd = SGD(lr=0.1)
model1_3.compile(loss='binary_crossentropy', optimizer='sgd')
model1_3.fit(X, y, batch_size=2, epochs=400)
print(model1_3.predict_proba(X).reshape(4*n))
scores1_3 = model1_3.evaluate(X, y)

# 2 Layers
model2_3 = Sequential()
model2_3.add(Dense(3, input_dim=2, activation='tanh')) #first Layer
model2_3.add(Dense(3, activation='tanh')) #second Layer
model2_3.add(Dense(1, activation='sigmoid'))
sgd = SGD(lr=0.1)
model2_3.compile(loss='binary_crossentropy', optimizer='sgd')
model2_3.fit(X, y, batch_size=2, epochs=400)
print(model2_3.predict_proba(X).reshape(4*n))
scores2_3 = model2_3.evaluate(X, y) # evaluate the model

# 3 Layers
model3_3 = Sequential()
model3_3.add(Dense(3, input_dim=2, activation='tanh')) #first Layer
model3_3.add(Dense(3, activation='tanh')) #second Layer
model3_3.add(Dense(3, activation='tanh')) #third Layer
model3_3.add(Dense(1, activation='sigmoid'))
sgd = SGD(lr=0.1)
model3_3.compile(loss='binary_crossentropy', optimizer='sgd')
model3_3.fit(X, y, batch_size=2, epochs=400)
print(model3_3.predict_proba(X).reshape(4*n))
scores3_3 = model3_3.evaluate(X, y) # evaluate the model

# 4 Layers
model4_3 = Sequential()
model4_3.add(Dense(3, input_dim=2, activation='tanh')) #first Layer
model4_3.add(Dense(3, activation='tanh')) #second Layer
model4_3.add(Dense(3, activation='tanh')) #third Layer
model4_3.add(Dense(3, activation='tanh')) #fourth Layer
model4_3.add(Dense(1, activation='sigmoid'))
sgd = SGD(lr=0.1)
model4_3.compile(loss='binary_crossentropy', optimizer='sgd')
model4_3.fit(X, y, batch_size=2, epochs=400)
print(model4_3.predict_proba(X).reshape(4*n))
scores4_3 = model4_3.evaluate(X, y) # evaluate the model

# 5 Layers
model5_3 = Sequential()
model5_3.add(Dense(3, input_dim=2, activation='tanh')) #first Layer
model5_3.add(Dense(3, activation='tanh')) #second Layer
model5_3.add(Dense(3, activation='tanh')) #third Layer
model5_3.add(Dense(3, activation='tanh')) #fourth Layer
model5_3.add(Dense(3, activation='tanh')) #fifth Layer
model5_3.add(Dense(1, activation='sigmoid'))
sgd = SGD(lr=0.1)
model5_3.compile(loss='binary_crossentropy', optimizer='sgd')

```

```

model5_3.fit(X, y, batch_size=2, epochs=400)
print(model5_3.predict_proba(X).reshape(4*n))
scores5_3 = model5_3.evaluate(X, y) # evaluate the model
0.9930178 0.8335678 0.9727447 0.9949606 0.9957796 0.7714665
0.9964323 0.9963763 0.94380397 0.9933458 0.99600005 0.98400676
0.99631226 0.99663216 0.9962824 0.99385124 0.9942194 0.9906001
0.9929278 0.9938419 0.99475217 0.9952544 0.99090946 0.99287117
0.9934825 0.99242556 0.01773462 0.00457767 0.00425583 0.00451839
0.00447074 0.00427061 0.00436348 0.01159886 0.00487137 0.00426149
0.00442171 0.00445125 0.00424767 0.02885684 0.00425577 0.8132191
0.00758037 0.00425449 0.00440139 0.25771248 0.004262 0.00433728
0.04112327 0.00464192 0.00436077 0.00457802 0.00425279 0.00427502
0.00427595 0.00430402 0.00452453 0.01799223 0.00441068 0.00431323
0.00424302 0.00440365 0.01283097 0.00426871 0.00459141 0.00425866
0.00450107 0.00763521 0.00663856 0.00755537 0.00688237 0.00705406
0.00713941 0.00488397 0.00890303 0.00705037 0.00725171 0.00713745
0.00624803 0.00497016 0.00544262 0.7473177 0.01698297 0.00681308
0.00680828 0.01970136 0.008091 0.00677314 0.10041916 0.008481
0.00678292 0.00427267 0.00831807 0.00761643 0.0073989 0.00677887
0.00482103 0.05042472 0.00726756 0.00684345 0.00636938 0.00693497
0.02555111 0.00607866 0.00820667 0.00586253]
5/5 [=====] - 0s 1ms/step - loss: 0.0331

```

```

In [19]: df_3 = np.array([scores1_3, scores2_3, scores3_3, scores4_3, scores5_3])
df_3

```

```

Out[19]: array([0.16349804, 0.09726368, 0.31269667, 0.02364234, 0.03312378])

```

4-neuron models

```

In [20]: # 1 Layer
model1_4 = Sequential()
model1_4.add(Dense(4, input_dim=2, activation='tanh')) #first Layer
model1_4.add(Dense(1, activation='sigmoid'))
sgd = SGD(lr=0.1)
model1_4.compile(loss='binary_crossentropy', optimizer='sgd')
model1_4.fit(X, y, batch_size=2, epochs=400)
print(model1_4.predict_proba(X).reshape(4*n))
scores1_4 = model1_4.evaluate(X, y)

# 2 Layers
model2_4 = Sequential()
model2_4.add(Dense(4, input_dim=2, activation='tanh')) #first Layer
model2_4.add(Dense(4, activation='tanh')) #second Layer
model2_4.add(Dense(1, activation='sigmoid'))
sgd = SGD(lr=0.1)
model2_4.compile(loss='binary_crossentropy', optimizer='sgd')
model2_4.fit(X, y, batch_size=2, epochs=400)
print(model2_4.predict_proba(X).reshape(4*n))
scores2_4 = model2_4.evaluate(X, y) # evaluate the model

# 3 Layers
model3_4 = Sequential()
model3_4.add(Dense(4, input_dim=2, activation='tanh')) #first Layer
model3_4.add(Dense(4, activation='tanh')) #second Layer
model3_4.add(Dense(4, activation='tanh')) #third Layer
model3_4.add(Dense(1, activation='sigmoid'))
sgd = SGD(lr=0.1)
model3_4.compile(loss='binary_crossentropy', optimizer='sgd')
model3_4.fit(X, y, batch_size=2, epochs=400)
print(model3_4.predict_proba(X).reshape(4*n))
scores3_4 = model3_4.evaluate(X, y) # evaluate the model

# 4 Layers
model4_4 = Sequential()
model4_4.add(Dense(4, input_dim=2, activation='tanh')) #first Layer
model4_4.add(Dense(4, activation='tanh')) #second Layer
model4_4.add(Dense(4, activation='tanh')) #third Layer
model4_4.add(Dense(4, activation='tanh')) #fourth Layer
model4_4.add(Dense(1, activation='sigmoid'))
sgd = SGD(lr=0.1)
model4_4.compile(loss='binary_crossentropy', optimizer='sgd')
model4_4.fit(X, y, batch_size=2, epochs=400)
print(model4_4.predict_proba(X).reshape(4*n))
scores4_4 = model4_4.evaluate(X, y) # evaluate the model

# 5 Layers
model5_4 = Sequential()
model5_4.add(Dense(4, input_dim=2, activation='tanh')) #first Layer
model5_4.add(Dense(4, activation='tanh')) #second Layer
model5_4.add(Dense(4, activation='tanh')) #third Layer
model5_4.add(Dense(4, activation='tanh')) #fourth Layer
model5_4.add(Dense(4, activation='tanh')) #fifth Layer
model5_4.add(Dense(1, activation='sigmoid'))
sgd = SGD(lr=0.1)
model5_4.compile(loss='binary_crossentropy', optimizer='sgd')

```

```

model5_4.fit(X, y, batch_size=2, epochs=400)
print(model5_4.predict_proba(X).reshape(4*n))
scores5_4 = model5_4.evaluate(X, y) # evaluate the model
9.9881852e-01 9.9642766e-01 9.9855804e-01 9.9859798e-01 9.9867213e-01
9.9854243e-01 9.9465132e-01 9.9863374e-01 9.9848831e-01 9.9864197e-01
9.4312429e-04 2.1353662e-03 1.0594726e-03 1.8286109e-03 3.0544996e-03
1.1482537e-03 1.2022555e-03 9.4985962e-04 1.0387003e-03 1.0983944e-03
1.5938282e-03 2.0766854e-03 1.0147691e-03 1.0327697e-03 9.6467137e-04
4.5830736e-01 1.3356507e-03 1.0603964e-03 2.5803149e-03 8.6001158e-03
1.1397600e-03 2.4036169e-03 5.7303011e-03 1.3760328e-03 2.3299456e-03
8.4057450e-04 1.0600984e-03 1.3822913e-03 1.2838244e-03 1.1356473e-03
8.8140368e-04 3.3524036e-03 1.0711551e-03 1.1507273e-03 9.8606944e-04
2.1381974e-03 1.1148363e-02 9.9623203e-04 1.4131963e-03 9.6538663e-04
2.1280646e-03 3.7652850e-03 2.2040904e-03 3.5057664e-03 3.7088990e-03
2.4352670e-03 2.6372671e-03 2.4102032e-03 2.9561222e-03 2.3311377e-03
3.1637549e-03 3.4661293e-03 2.0642281e-03 2.4262369e-03 1.9360185e-03
6.2009609e-01 7.9893768e-03 2.2216439e-03 3.4881830e-03 1.4453262e-02
2.4994612e-03 3.3524334e-03 7.3363870e-02 3.3100843e-03 3.3559799e-03
1.7821193e-03 2.3247898e-03 2.8271079e-03 2.6907921e-03 2.4204850e-03
1.7154813e-03 1.7299235e-02 2.4747849e-03 2.4637878e-03 2.0032823e-03
3.3796132e-03 1.9443780e-02 2.1184087e-03 3.2892823e-03 2.0761788e-03]
5/5 [=====] - 0s 1ms/step - loss: 0.0207

```

```

In [21]: df_4 = np.array([scores1_4, scores2_4, scores3_4, scores4_4, scores5_4])
df_4

```

```

Out[21]: array([0.15961495, 0.03270616, 0.02293862, 0.02196405, 0.02074153])

```

Part 1 - continued

- Using the most optimal configuraion (n-layers, k-neurons per layer), compare how tanh , sigmoid , softplus and relu effect the loss after 400 epochs. Try other Activation functions as well (<https://keras.io/activations/>)

```

In [28]: df

```

```

Out[28]: array([0.31730145, 0.31588522, 0.36067224, 0.31026211, 0.0482264 ])

```

```

In [29]: df_3

```

```

Out[29]: array([0.16349804, 0.09726368, 0.31269667, 0.02364234, 0.03312378])

```

```

In [30]: df_4

```

```

Out[30]: array([0.15961495, 0.03270616, 0.02293862, 0.02196405, 0.02074153])

```

The model with 5 layers and 4 neurons resulted in lowest loss. Will use to compare with other models

```

In [32]: model5_4 = Sequential()
model5_4.add(Dense(4, input_dim=2, activation='tanh')) #first layer
model5_4.add(Dense(4, activation='tanh')) #second layer
model5_4.add(Dense(4, activation='tanh')) #third layer
model5_4.add(Dense(4, activation='tanh')) #fourth layer
model5_4.add(Dense(4, activation='tanh')) #fifth layer
model5_4.add(Dense(1, activation='sigmoid'))
sgd = SGD(lr=0.1)
model5_4.compile(loss='binary_crossentropy', optimizer='sgd')
model5_4.fit(X, y, batch_size=2, epochs=400)
print(model5_4.predict_proba(X).reshape(4*n))
scores_tanh = model5_4.evaluate(X, y) # evaluate the model

# sigmoid
model_sig = Sequential()
model_sig.add(Dense(4, input_dim=2, activation='sigmoid')) #first layer
model_sig.add(Dense(4, activation='sigmoid')) #second layer
model_sig.add(Dense(4, activation='sigmoid')) #third layer
model_sig.add(Dense(4, activation='sigmoid')) #fourth layer
model_sig.add(Dense(4, activation='sigmoid')) #fifth layer
model_sig.add(Dense(1, activation='sigmoid'))
sgd = SGD(lr=0.1)
model_sig.compile(loss='binary_crossentropy', optimizer='sgd')
model_sig.fit(X, y, batch_size=2, epochs=400)
print(model_sig.predict_proba(X).reshape(4*n))
scores_sig = model_sig.evaluate(X, y) # evaluate the model

# softplus
model_soft = Sequential()
model_soft.add(Dense(4, input_dim=2, activation='softplus')) #first layer
model_soft.add(Dense(4, activation='softplus')) #second layer
model_soft.add(Dense(4, activation='softplus')) #third layer
model_soft.add(Dense(4, activation='softplus')) #fourth layer
model_soft.add(Dense(4, activation='softplus')) #fifth layer
model_soft.add(Dense(1, activation='sigmoid'))
sgd = SGD(lr=0.1)
model_soft.compile(loss='binary_crossentropy', optimizer='sgd')
model_soft.fit(X, y, batch_size=2, epochs=400)
print(model_soft.predict_proba(X).reshape(4*n))
scores_soft = model_soft.evaluate(X, y) # evaluate the model

# relu
model_relu = Sequential()
model_relu.add(Dense(4, input_dim=2, activation='relu')) #first layer
model_relu.add(Dense(4, activation='relu')) #second layer
model_relu.add(Dense(4, activation='relu')) #third layer
model_relu.add(Dense(4, activation='relu')) #fourth layer
model_relu.add(Dense(4, activation='relu')) #fifth layer
model_relu.add(Dense(1, activation='sigmoid'))
sgd = SGD(lr=0.1)
model_relu.compile(loss='binary_crossentropy', optimizer='sgd')
model_relu.fit(X, y, batch_size=2, epochs=400)
print(model_relu.predict_proba(X).reshape(4*n))
scores_relu = model_relu.evaluate(X, y) # evaluate the model

```

```

80/80 [=====] - 0s 1ms/step - loss: 0.6654
Epoch 11/400

```

```
Epoch 11/400  
80/80 [=====] - 0s 1ms/step - loss: 0.6584  
Epoch 12/400  
80/80 [=====] - 0s 935us/step - loss: 0.6519  
Epoch 13/400  
80/80 [=====] - 0s 1ms/step - loss: 0.6437  
Epoch 14/400  
80/80 [=====] - 0s 997us/step - loss: 0.6380  
Epoch 15/400  
80/80 [=====] - 0s 972us/step - loss: 0.6311  
Epoch 16/400  
80/80 [=====] - 0s 1ms/step - loss: 0.6251  
Epoch 17/400  
80/80 [=====] - 0s 985us/step - loss: 0.6204  
Epoch 18/400  
80/80 [=====] - 0s 947us/step - loss: 0.6154  
Epoch 19/400  
80/80 [=====] - 0s 1ms/step - loss: 0.6114  
Epoch 20/400
```

In [33]: scores_tanh

Out[33]: 0.020530756562948227

In [34]: scores_sig

Out[34]: 0.6931532621383667

In [36]: scores_soft

Out[36]: 0.0738680362701416

In [37]: scores_relu

Out[37]: 0.022404631599783897

Part 1 - continued

4. Instead of SGD try other optimizers and report on the loss score. (<https://keras.io/optimizers/>)

will use the tanh model with SGD and compare to other optimizers

```

In [40]: # adam optimizer
from tensorflow import keras
from tensorflow.keras import layers

model_adam = Sequential()
model_adam.add(Dense(4, input_dim=2, activation='tanh')) #first layer
model_adam.add(Dense(4, activation='tanh')) #second layer
model_adam.add(Dense(4, activation='tanh')) #third layer
model_adam.add(Dense(4, activation='tanh')) #fourth layer
model_adam.add(Dense(4, activation='tanh')) #fifth layer
model_adam.add(Dense(1, activation='sigmoid'))
adam = keras.optimizers.Adam(learning_rate=0.01)
model_adam.compile(loss='binary_crossentropy', optimizer='adam')
model_adam.fit(X, y, batch_size=2, epochs=400)
print(model_adam.predict_proba(X).reshape(4*n))
scores_adam = model_adam.evaluate(X, y) # evaluate the model


0.9904478e-01 0.9630034e-01 0.9829805e-01 0.9843597e-01 0.9879003e-01
9.9820125e-01 9.9532622e-01 9.9849010e-01 9.9807596e-01 9.9849379e-01
6.7308545e-04 2.2098422e-04 3.5175681e-04 2.3555756e-04 2.9137731e-04
3.6633015e-04 2.7906895e-04 2.9352307e-04 1.6680360e-04 3.6087632e-04
2.7540326e-04 2.7349591e-04 3.4230947e-04 5.1289797e-04 2.9397011e-04
6.9679499e-01 5.6260824e-04 3.5387278e-04 3.3193827e-04 6.6405833e-03
3.6600232e-04 3.8492680e-04 7.0707500e-03 2.0051003e-04 3.6174059e-04
2.7599931e-04 3.5184622e-04 3.8340688e-04 3.7926435e-04 3.1465292e-04
3.2043457e-04 3.6230683e-04 2.1570921e-04 3.0982494e-04 3.3518672e-04
3.1182170e-04 4.6333671e-04 2.4867058e-04 2.0974874e-04 2.2357702e-04
8.1494451e-04 1.4898777e-03 1.2208223e-03 1.4412105e-03 1.3707280e-03
1.2377501e-03 1.3180673e-03 1.0784864e-03 1.8896461e-03 1.2298226e-03
1.3602674e-03 1.3792813e-03 1.2030005e-03 8.5097551e-04 1.1889935e-03
6.0804570e-01 5.0607920e-03 1.2210608e-03 1.3224185e-03 5.6765676e-03
1.2376010e-03 1.2777746e-03 2.5666893e-02 1.5205145e-03 1.2945831e-03
9.9918246e-04 1.2273490e-03 1.2497008e-03 1.2467206e-03 1.2682676e-03
9.9945068e-04 5.7137609e-03 1.3881326e-03 1.2759566e-03 1.1903048e-03
1.3345480e-03 6.1666965e-03 1.2716353e-03 1.4884174e-03 1.2865663e-03]
5/5 [=====] - 0s 997us/step - loss: 0.0194

```



```
In [45]: plt.scatter(*zip(*X), c=model_adam.predict(X).reshape(4*n))
```

```
Out[45]: <matplotlib.collections.PathCollection at 0x170c29d6348>
```



Part 2 - BYOD (Bring your own Dataset)

Using your own dataset, experiment and find the best Neural Network configuration. You may use any resource to improve results, just reference it.

While you may use any dataset, I'd prefer you didn't use the diabetes dataset used in the lesson.

<https://stackoverflow.com/questions/34673164/how-to-train-and-tune-an-artificial-multilayer-perceptron-neural-network-using-k> (<https://stackoverflow.com/questions/34673164/how-to-train-and-tune-an-artificial-multilayer-perceptron-neural-network-using-k>)

<https://keras.io/> (<https://keras.io/>)

Using Heart Disease data

<https://www.kaggle.com/volodymyrgavrysh/heart-disease?select=heart.csv>
(<https://www.kaggle.com/volodymyrgavrysh/heart-disease?select=heart.csv>)

Using 13 factors to predict diagnosis of heart disease

Attribute Information:

Age: Age

Sex: Sex (1 = male; 0 = female)

ChestPain: Chest pain (typical, asymptotic, nonanginal, nontypical)

RestBP: Resting blood pressure

Chol: Serum cholestoral in mg/dl

Fbs: Fasting blood sugar > 120 mg/dl (1 = true; 0 = false)
 RestECG: Resting electrocardiographic results
 MaxHR: Maximum heart rate achieved
 ExAng: Exercise induced angina (1 = yes; 0 = no)
 Oldpeak: ST depression induced by exercise relative to rest
 Slope: Slope of the peak exercise ST segment
 Ca: Number of major vessels colored by flourosopy (0 - 3)
 Thal: (3 = normal; 6 = fixed defect; 7 = reversable defect)
 target: AHD - Diagnosis of heart disease (1 = yes; 0 = no)

```
In [178]: # Load dataset
dataset = pd.read_csv('../data/heart2.csv', index_col=False)
dataset
```

```
Out[178]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1
...
298	57	0	0	140	241	0	1	123	1	0.2	1	0	3	0
299	45	1	3	110	264	0	1	132	0	1.2	1	0	3	0
300	68	1	0	144	193	1	1	141	0	3.4	1	2	3	0
301	57	1	0	130	131	0	1	115	1	1.2	1	1	3	0
302	57	0	1	130	236	0	0	174	0	0.0	1	1	2	0

303 rows × 14 columns

```
In [179]: # split into input (X) and output (Y) variables
X = dataset.iloc[:,0:13]
Y = dataset.iloc[:,13]
```

```

In [183]: # create models with 2, 4, and 6 layers
model = Sequential()
model.add(Dense(24, input_dim=13, activation='tanh'))
model.add(Dense(24, activation='tanh'))
model.add(Dense(1, activation='sigmoid'))
# Compile model
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
# Fit the model
model.fit(X, Y, epochs=400, batch_size=10)
# evaluate the model
scores2 = model.evaluate(X, Y)
print("\ns: %.2f%%" % (model.metrics_names[1], scores[1]*100))

# 4 layers
model = Sequential()
model.add(Dense(24, input_dim=13, activation='tanh'))
model.add(Dense(24, activation='tanh'))
model.add(Dense(24, activation='tanh'))
model.add(Dense(24, activation='tanh'))
model.add(Dense(1, activation='sigmoid'))
# Compile model
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
# Fit the model
model.fit(X, Y, epochs=400, batch_size=10)
# evaluate the model
scores4 = model.evaluate(X, Y)
print("\ns: %.2f%%" % (model.metrics_names[1], scores[1]*100))

# 6 layers
model = Sequential()
model.add(Dense(24, input_dim=13, activation='tanh'))
model.add(Dense(24, activation='tanh'))
model.add(Dense(24, activation='tanh'))
model.add(Dense(24, activation='tanh'))
model.add(Dense(24, activation='tanh'))
model.add(Dense(24, activation='tanh'))
model.add(Dense(1, activation='sigmoid'))
# Compile model
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
# Fit the model
model.fit(X, Y, epochs=400, batch_size=10)
# evaluate the model
scores6 = model.evaluate(X, Y)
print("\ns: %.2f%%" % (model.metrics_names[1], scores[1]*100))

```

y: 0.7129

Epoch 174/400

31/31 [=====] - 0s 2ms/step - loss: 0.5339 - accurac

y: 0.7327

Epoch 175/400

31/31 [=====] - 0s 2ms/step - loss: 0.5175 - accurac

y: 0.7294

Epoch 176/400

31/31 [=====] - 0s 2ms/step - loss: 0.5319 - accurac

y: 0.7162

Epoch 177/400

```
31/31 [=====] - 0s 2ms/step - loss: 0.5366 - accurac
y: 0.6898
Epoch 178/400
31/31 [=====] - ETA: 0s - loss: 0.5023 - accuracy:
0.80 - 0s 2ms/step - loss: 0.5314 - accuracy: 0.7195
Epoch 179/400
31/31 [=====] - 0s 2ms/step - loss: 0.5432 - accurac
y: 0.7459
Epoch 180/400
```

In [184]: scores2

Out[184]: [0.3897251784801483, 0.7788779139518738]

In [185]: scores4

Out[185]: [0.30145254731178284, 0.8811880946159363]

In [186]: scores6

Out[186]: [0.38527220487594604, 0.8448845148086548]

4 layer model had lowest loss score. Will explore different number of neurons

In [187]:

```
# 4 layers, 36 neurons
model = Sequential()
model.add(Dense(36, input_dim=13, activation='tanh'))
model.add(Dense(36, activation='tanh'))
model.add(Dense(36, activation='tanh'))
model.add(Dense(36, activation='tanh'))
model.add(Dense(1, activation='sigmoid'))
# Compile model
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
# Fit the model
model.fit(X, Y, epochs=400, batch_size=10)
# evaluate the model
scores4_36 = model.evaluate(X, Y)
print("\ns: %.2f%%" % (model.metrics_names[1], scores[1]*100))

# 4 layers, 48 neurons
model = Sequential()
model.add(Dense(48, input_dim=13, activation='tanh'))
model.add(Dense(48, activation='tanh'))
model.add(Dense(48, activation='tanh'))
model.add(Dense(48, activation='tanh'))
model.add(Dense(1, activation='sigmoid'))
# Compile model
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
# Fit the model
model.fit(X, Y, epochs=400, batch_size=10)
# evaluate the model
scores4_48 = model.evaluate(X, Y)
print("\ns: %.2f%%" % (model.metrics_names[1], scores[1]*100))
```

```
Epoch 251/400
31/31 [=====] - 0s 2ms/step - loss: 0.3401 - accuracy: 0.8482
Epoch 252/400
31/31 [=====] - 0s 1ms/step - loss: 0.3601 - accuracy: 0.8284
Epoch 253/400
31/31 [=====] - 0s 2ms/step - loss: 0.3693 - accuracy: 0.8548
Epoch 254/400
31/31 [=====] - 0s 1ms/step - loss: 0.3539 - accuracy: 0.8449
Epoch 255/400
31/31 [=====] - 0s 2ms/step - loss: 0.3589 - accuracy: 0.8350
Epoch 256/400
31/31 [=====] - 0s 1ms/step - loss: 0.3515 - accuracy: 0.8548
Epoch 257/400
31/31 [=====] - 0s 2ms/step - loss: 0.3400 - accuracy:
```

```
In [188]: scores4
```

```
Out[188]: [0.30145254731178284, 0.8811880946159363]
```

```
In [189]: scores4_36
```

```
Out[189]: [0.27431437373161316, 0.8415841460227966]
```

```
In [190]: scores4_48
```

```
Out[190]: [0.3702474534511566, 0.8646864891052246]
```

4 layer with 36 neurons still has lowest loss score. Will adjust activation

```

In [191]: # 4 layers, 36 neurons, sigmoid
model = Sequential()
model.add(Dense(36, input_dim=13, activation='sigmoid'))
model.add(Dense(36, activation='sigmoid'))
model.add(Dense(36, activation='sigmoid'))
model.add(Dense(36, activation='sigmoid'))
model.add(Dense(1, activation='sigmoid'))
# Compile model
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
# Fit the model
model.fit(X, Y, epochs=400, batch_size=10)
# evaluate the model
scores4_36_sig = model.evaluate(X, Y)
print("\ns: %.2f%%" % (model.metrics_names[1], scores[1]*100))

# 4 layers, 36 neurons, softplus
model = Sequential()
model.add(Dense(36, input_dim=13, activation='softplus'))
model.add(Dense(36, activation='softplus'))
model.add(Dense(36, activation='softplus'))
model.add(Dense(36, activation='softplus'))
model.add(Dense(1, activation='sigmoid'))
# Compile model
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
# Fit the model
model.fit(X, Y, epochs=400, batch_size=10)
# evaluate the model
scores4_36_soft = model.evaluate(X, Y)
print("\ns: %.2f%%" % (model.metrics_names[1], scores[1]*100))

# 4 layers, 36 neurons, relu
model = Sequential()
model.add(Dense(36, input_dim=13, activation='relu'))
model.add(Dense(36, activation='relu'))
model.add(Dense(36, activation='relu'))
model.add(Dense(36, activation='relu'))
model.add(Dense(1, activation='sigmoid'))
# Compile model
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
# Fit the model
model.fit(X, Y, epochs=400, batch_size=10)
# evaluate the model
scores4_36_relu = model.evaluate(X, Y)
print("\ns: %.2f%%" % (model.metrics_names[1], scores[1]*100))

```

```

31/31 [-----] - 0s 1ms/step - loss: 0.5548 - accuracy: 0.6964
Epoch 18/400
31/31 [=====] - 0s 2ms/step - loss: 0.5727 - accuracy: 0.7294
Epoch 19/400
31/31 [=====] - 0s 1ms/step - loss: 0.5689 - accuracy: 0.7195
Epoch 20/400
31/31 [=====] - 0s 2ms/step - loss: 0.5574 - accuracy: 0.7228
Epoch 21/400
31/31 [=====] - 0s 2ms/step - loss: 0.5530 - accuracy: 0.7228

```

```
31/31 [=====] - 0s 2ms/step - loss: 0.5528 - accurac
y: 0.7261
Epoch 22/400
31/31 [=====] - 0s 1ms/step - loss: 0.5480 - accurac
y: 0.7261
Epoch 23/400
31/31 [=====] - 0s 2ms/step - loss: 0.5377 - accurac
y: 0.7294
```

In [192]: scores4_36

Out[192]: [0.27431437373161316, 0.8415841460227966]

In [193]: scores4_36_sig

Out[193]: [0.2899039089679718, 0.8745874762535095]

In [194]: scores4_36_soft

Out[194]: [0.0405021533370018, 0.9900990128517151]

In [195]: scores4_36_relu

Out[195]: [0.15602846443653107, 0.933993399143219]

Conclusion:

The model with 4 layers, 36 neurons, and activation with `softplus`, was the best performing model with 0.0405 loss and 0.99 accuracy