

Video 13.1 <https://www.youtube.com/watch?v=kIGHE7Cfe1s> (<https://www.youtube.com/watch?v=kIGHE7Cfe1s>)

Video 13.2 <https://www.youtube.com/watch?v=Rm9bJcDd1KU> (<https://www.youtube.com/watch?v=Rm9bJcDd1KU>)

Video 13.3 <https://youtu.be/6HjZk-3LsjE> (<https://youtu.be/6HjZk-3LsjE>)

## Assignment

1. change the `encoding_dim` through various values ( `range(2,18,2)` ) and store or keep track of the best loss you can get. Plot the 8 pairs of dimensions vs loss on a scatter plot

```
In [1]: from keras.callbacks import TensorBoard
from keras.callbacks import EarlyStopping
from keras.layers import Input, Dense
from keras.models import Model
from keras.datasets import mnist
import pandas as pd
import numpy as np

(xtrain, ytrain), (xtest, ytest) = mnist.load_data()

xtrain = xtrain.astype('float32') / 255.
xtest = xtest.astype('float32') / 255.
xtrain = xtrain.reshape((len(xtrain), np.prod(xtrain.shape[1:])))
xtest = xtest.reshape((len(xtest), np.prod(xtest.shape[1:])))
xtrain.shape, xtest.shape
```

Using TensorFlow backend.

```
Out[1]: ((60000, 784), (10000, 784))
```

```
In [5]: import tensorflow as tf
from tensorflow import keras
```

Documentation for writing callbacks found at: [https://keras.io/guides/writing\\_your\\_own\\_callbacks/](https://keras.io/guides/writing_your_own_callbacks/) ([https://keras.io/guides/writing\\_your\\_own\\_callbacks/](https://keras.io/guides/writing_your_own_callbacks/))

Documentation for Earlystop at: <https://keras.io/api/callbacks/#earlystopping> (<https://keras.io/api/callbacks/#earlystopping>)

```

In [6]: loss = {}

for i in range(2, 18, 2):

    encoding_dim = i

    x = input_img = Input(shape=(784,))
    x = Dense(256, activation='relu')(x)
    x = Dense(128, activation='relu')(x)
    encoded = Dense(encoding_dim, activation='relu')(x)

    x = Dense(128, activation='relu')(encoded)
    x = Dense(256, activation='relu')(x)
    decoded = Dense(784, activation='sigmoid')(x)

    autoencoder = Model(input_img, decoded)

    encoder = Model(input_img, encoded)

    encoded_input = Input(shape=(encoding_dim,))

    dcd1 = autoencoder.layers[-1]
    dcd2 = autoencoder.layers[-2]
    dcd3 = autoencoder.layers[-3]

    decoder = Model(encoded_input, dcd1(dcd2(dcd3(encoded_input))))

    autoencoder.compile(optimizer='adadelata', loss='binary_crossentropy')

    autoencoder.fit(xtrain, xtrain,
                    epochs=50,
                    batch_size=256,
                    shuffle=True,
                    validation_data=(xtest, xtest),
                    callbacks=[tf.keras.callbacks.EarlyStopping(patience=2)])

    loss[i] = autoencoder.evaluate(xtrain, xtrain, verbose = 0)

```


In [20]: loss

```
Out[20]: {2: 0.6900376677513123,
          4: 0.6899483799934387,
          6: 0.6895860433578491,
          8: 0.6895838975906372,
          10: 0.6892221570014954,
          12: 0.6897948384284973,
          14: 0.6827694773674011,
          16: 0.6865726113319397}
```

```
In [21]: import matplotlib.pyplot as plt
         %matplotlib inline
```

```
In [27]: plt.scatter(loss.keys(), loss.values())
         plt.xlabel("Dimensions")
         plt.ylabel('Loss')
```

Out[27]: Text(0, 0.5, 'Loss')



**2. using the previous assignment's model of detecting images, how does the accuracy change when you run the digit-prediction model on these 'decoded' values?**

```
In [28]: import keras
         from keras.datasets import mnist
         from keras.models import Sequential
         from keras.optimizers import RMSprop
         from keras.layers import Dense, Dropout, Flatten
         from keras.layers import Conv2D, MaxPooling2D
         from keras import backend
```

```
In [29]: (xtrain, ytrain), (xtest, ytest) = mnist.load_data()

xtrain = xtrain.astype('float32') / 255.
xtest = xtest.astype('float32') / 255.
xtrain = xtrain.reshape((len(xtrain), np.prod(xtrain.shape[1:])))
xtest = xtest.reshape((len(xtest), np.prod(xtest.shape[1:])))
xtrain.shape, xtest.shape
```

```
Out[29]: ((60000, 784), (10000, 784))
```

```
In [31]: batch_size = 128
num_classes = 10
epochs = 20


# convert class vectors to binary class matrices
ytrain = keras.utils.to_categorical(ytrain, num_classes)
ytest = keras.utils.to_categorical(ytest, num_classes)
```

```
In [33]: import matplotlib.pyplot as plt
```

```
In [34]: encoded_imgs = encoder.predict(xtest)
decoded_imgs = decoder.predict(encoded_imgs)

n = 20 # how many digits we will display
plt.figure(figsize=(40, 4))
for i in range(n):
    # display original
    ax = plt.subplot(2, n, i + 1)
    plt.imshow(xtest[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)

    # display reconstruction
    ax = plt.subplot(2, n, i + 1 + n)
    plt.imshow(decoded_imgs[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
plt.show()
```



```
In [35]: model = Sequential()
model.add(Dense(512, activation='relu', input_shape=(784,)))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(10, activation='softmax'))

model.summary()

model.compile(loss='categorical_crossentropy',
              optimizer=RMSprop(),
              metrics=['accuracy'])

history = model.fit(xtrain, ytrain,
                   batch_size=batch_size,
                   epochs=epochs,
                   verbose=1,
                   validation_data=(decoded_imgs, ytest))
score_nn = model.evaluate(decoded_imgs, ytest, verbose=0)
print('Test loss:', score_nn[0])
print('Test accuracy:', score_nn[1])
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
dense_81 (Dense)	(None, 512)	401920
dropout_2 (Dropout)	(None, 512)	0
dense_82 (Dense)	(None, 512)	262656
dropout_3 (Dropout)	(None, 512)	0
dense_83 (Dense)	(None, 10)	5130
Total params: 669,706		
Trainable params: 669,706		
Non-trainable params: 0		

```
Epoch 1/20
469/469 [=====] - 8s 18ms/step - loss: 0.2450 - accuracy: 0.9237 - val_loss: 5.5860 - val_accuracy: 0.0892
Epoch 2/20
469/469 [=====] - 8s 16ms/step - loss: 0.1016 - accuracy: 0.9689 - val_loss: 10.6933 - val_accuracy: 0.0892
Epoch 3/20
469/469 [=====] - 8s 16ms/step - loss: 0.0756 - accuracy: 0.9779 - val_loss: 15.5065 - val_accuracy: 0.0958
Epoch 4/20
469/469 [=====] - 8s 17ms/step - loss: 0.0608 - accuracy: 0.9817 - val_loss: 17.0012 - val_accuracy: 0.0958
Epoch 5/20
469/469 [=====] - 9s 19ms/step - loss: 0.0499 - accuracy: 0.9846 - val_loss: 27.6631 - val_accuracy: 0.0958
Epoch 6/20
```

```

469/469 [=====] - 10s 20ms/step - loss: 0.0433 - accuracy: 0.9874 - val_loss: 29.8694 - val_accuracy: 0.1032
Epoch 7/20
469/469 [=====] - 8s 18ms/step - loss: 0.0371 - accuracy: 0.9886 - val_loss: 43.2751 - val_accuracy: 0.1032
Epoch 8/20
469/469 [=====] - 9s 20ms/step - loss: 0.0338 - accuracy: 0.9896 - val_loss: 46.0997 - val_accuracy: 0.0958
Epoch 9/20
469/469 [=====] - 8s 17ms/step - loss: 0.0311 - accuracy: 0.9911 - val_loss: 69.4659 - val_accuracy: 0.1032
Epoch 10/20
469/469 [=====] - 8s 17ms/step - loss: 0.0281 - accuracy: 0.9913 - val_loss: 61.7929 - val_accuracy: 0.1032
Epoch 11/20
469/469 [=====] - 8s 17ms/step - loss: 0.0265 - accuracy: 0.9923 - val_loss: 54.4139 - val_accuracy: 0.0958
Epoch 12/20
469/469 [=====] - 8s 17ms/step - loss: 0.0244 - accuracy: 0.9926 - val_loss: 63.5142 - val_accuracy: 0.1032
Epoch 13/20
469/469 [=====] - 8s 17ms/step - loss: 0.0228 - accuracy: 0.9935 - val_loss: 50.2890 - val_accuracy: 0.1032
Epoch 14/20
469/469 [=====] - 8s 17ms/step - loss: 0.0202 - accuracy: 0.9937 - val_loss: 75.2488 - val_accuracy: 0.0958
Epoch 15/20
469/469 [=====] - 8s 17ms/step - loss: 0.0198 - accuracy: 0.9944 - val_loss: 79.1382 - val_accuracy: 0.0958
Epoch 16/20
469/469 [=====] - 8s 17ms/step - loss: 0.0210 - accuracy: 0.9941 - val_loss: 85.6847 - val_accuracy: 0.0961
Epoch 17/20
469/469 [=====] - 8s 17ms/step - loss: 0.0193 - accuracy: 0.9949 - val_loss: 97.1585 - val_accuracy: 0.0958
Epoch 18/20
469/469 [=====] - 8s 17ms/step - loss: 0.0186 - accuracy: 0.9947 - val_loss: 109.5125 - val_accuracy: 0.1032
Epoch 19/20
469/469 [=====] - 8s 17ms/step - loss: 0.0176 - accuracy: 0.9956 - val_loss: 93.1350 - val_accuracy: 0.1032
Epoch 20/20
469/469 [=====] - 8s 17ms/step - loss: 0.0166 - accuracy: 0.9956 - val_loss: 90.6213 - val_accuracy: 0.0982
Test loss: 90.62125396728516
Test accuracy: 0.098200008225441

```

**3. apply noise to *only* the input of the autoencoder (not the output). demonstrate that your autoencoder can strip out noise.**

```

In [37]: xtrain_noise10 = xtrain + np.random.normal(0, 255*.10, xtrain.shape)
         xtest_noise10 = xtest + np.random.normal(0, 255*.10, xtest.shape)

```

```
In [39]: autoencoder.fit(xtrain_noise10, xtrain,
                        epochs=20,
                        batch_size=256,
                        shuffle=True,
                        validation_data=(xtest_noise10, xtest),
                        callbacks=[tf.keras.callbacks.EarlyStopping(patience=2)])
```

```
Epoch 1/20
235/235 [=====] - 7s 28ms/step - loss: 0.8259 - val_loss: 0.7846
Epoch 2/20
235/235 [=====] - 6s 26ms/step - loss: 0.7538 - val_loss: 0.7271
Epoch 3/20
235/235 [=====] - 6s 26ms/step - loss: 0.7066 - val_loss: 0.6887
Epoch 4/20
235/235 [=====] - 7s 30ms/step - loss: 0.6740 - val_loss: 0.6608
Epoch 5/20
235/235 [=====] - 8s 32ms/step - loss: 0.6488 - val_loss: 0.6375
Epoch 6/20
235/235 [=====] - 7s 32ms/step - loss: 0.6259 - val_loss: 0.6144
Epoch 7/20
235/235 [=====] - 8s 34ms/step - loss: 0.6015 - val_loss: 0.5884
Epoch 8/20
235/235 [=====] - 7s 30ms/step - loss: 0.5734 - val_loss: 0.5579
Epoch 9/20
235/235 [=====] - 7s 28ms/step - loss: 0.5406 - val_loss: 0.5227
Epoch 10/20
235/235 [=====] - 8s 32ms/step - loss: 0.5036 - val_loss: 0.4841
Epoch 11/20
235/235 [=====] - 7s 30ms/step - loss: 0.4645 - val_loss: 0.4450
Epoch 12/20
235/235 [=====] - 7s 30ms/step - loss: 0.4266 - val_loss: 0.4088
Epoch 13/20
235/235 [=====] - 7s 28ms/step - loss: 0.3931 - val_loss: 0.3784
Epoch 14/20
235/235 [=====] - 6s 27ms/step - loss: 0.3660 - val_loss: 0.3547
Epoch 15/20
235/235 [=====] - 7s 30ms/step - loss: 0.3455 - val_loss: 0.3373
Epoch 16/20
235/235 [=====] - 8s 32ms/step - loss: 0.3306 - val_loss: 0.3246
Epoch 17/20
```

```

235/235 [=====] - 7s 32ms/step - loss: 0.3197 - val_loss: 0.3153
Epoch 18/20
235/235 [=====] - 7s 29ms/step - loss: 0.3118 - val_loss: 0.3085
Epoch 19/20
235/235 [=====] - 7s 30ms/step - loss: 0.3059 - val_loss: 0.3033
Epoch 20/20
235/235 [=====] - 7s 29ms/step - loss: 0.3014 - val_loss: 0.2993

```

Out[39]: <tensorflow.python.keras.callbacks.History at 0x19b02bf8e08>

```

In [40]: encoded_imgs = encoder.predict(xtest_noise10)
         decoded_imgs = decoder.predict(encoded_imgs)

         n = 20 # how many digits we will display
         plt.figure(figsize=(40, 4))
         for i in range(n):
             # display original
             ax = plt.subplot(2, n, i + 1)
             plt.imshow(xtest[i].reshape(28, 28))
             plt.gray()
             ax.get_xaxis().set_visible(False)
             ax.get_yaxis().set_visible(False)

             # display reconstruction
             ax = plt.subplot(2, n, i + 1 + n)
             plt.imshow(decoded_imgs[i].reshape(28, 28))
             plt.gray()
             ax.get_xaxis().set_visible(False)
             ax.get_yaxis().set_visible(False)
         plt.show()

```

